

Application and Architecture Description

This application connects to Twitter and downloads tweets, parses them into their constituent words, counts the number of times each word has appeared and writes that information to a data table in Postgres. It does this in several steps, first it connects to the Twitter streaming API with three spouts collecting tweets. It then has two stages of bolts. In the first stage tweets are parsed into individual words, filtering out the contents of tweets that are not words like hashtags, retweets, mentions and URLs. These bolts then emit all valid words. The valid words are passed to the count bolts, which count the number of times each word has appeared and write that information to a database in Postgres. If a word is not in Postgres a new row is created for the word, if the word is in the database its count is incremented by 1.

File Dependencies

There are several file dependencies in the architecture. The `tweetwordcount.clj` file which controls the running of the spouts and bolts requires the `tweets.py` spout and the `parse.py` and `wordcount.py` bolts to run and those files have to be in the correct directory for the `tweetwordcount.clj` file to be able to find them.

The script is also dependent on a database named `tcount` and a data table named `tweetwordcount` already existing in Postgres.

Running the Application

To run the application make your current working directory `w205_2017_spring/exercise_2/exttweetwordcount` and type “sparse run” into the terminal. The process will run, collecting and parsing tweets, then writing the words and their counts to the data table, until it is killed by the user.

To run the two serving scripts change your directory to `/w205_2017_spring/exercise_2/exttweetwordcount/serving_scripts`. The two serving scripts are there. Each program should be run like any other Python program, with arguments after the program name.

File Structure

The git repo is called `w205_2017_spring` and all application information is in the `exercise_2` directory. In that directory there is the `readme.txt` which describes how to run the application, `plog.png` which has the bar chart, a `screenshots` directory which has screenshots to show various stages of the application working, a python file that stores my twitter credentials and an `exttweetwordcount` directory which contains all the program files. There is also a folder called `serving_scripts` which has the two Python serving scripts in it.

The extweetwordcount directory is substantially the same as the version I downloaded from the class Git repo. The topologies directory contains the modified application topology and the src directory contains the modified bolt and spout scripts. Files for each bolt and spout are in their respective directories. The rest of the files in the directory are either what was in the original repo or log files automatically created by the application in the logs directory.